



# Hashing the Hypertrie: Space- and Time-Efficient Indexing for SPARQL in Tensors

Anonymous Authors

Baseline hypertrie



Optimized hypertrie



**Figure 1: Exemplary visualization of the three proposed optimizations of the hypertrie. For better readability, only nodes relevant for the example are shown.**

## ABSTRACT

Time-efficient solutions for querying RDF knowledge graphs depend on indexing structures with low response times to answer SPARQL queries rapidly. Hypertries—an indexing structure recently developed for tensor-based triple stores—have achieved significant runtime improvements over several mainstream storage solutions for RDF knowledge graphs. However, the space footprint of this novel data structure is still often larger than that of many mainstream solutions. In this work, we show how to reduce the memory footprint of hypertries and thereby further speed up query processing in hypertrie-based RDF storage solutions. Our approach relies on three strategies: (1) the elimination of duplicates via hashing, (2) the compression of non-branching paths, and (3) the storage of single-entry leaf nodes in their parent nodes. We evaluate these strategies by comparing them with baseline hypertries as well as popular triple stores such as Virtuoso, Fuseki, GraphDB, Blazegraph and gStore. We rely on four datasets/benchmark generators in our evaluation: SWDF, DBpedia, WatDiv, and WikiData. Our results suggest that our modifications significantly reduce the memory footprint of hypertries by up to 70% while leading to a relative improvement of up to 39% with respect to average Queries per Second and up to 740% with respect to Query Mixes per Hour.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from [permissions@acm.org](mailto:permissions@acm.org).

Woodstock '18, June 03–05, 2018, Woodstock, NY

© 2018 Association for Computing Machinery.

ACM ISBN 978-1-4503-XXXX-X/18/06...\$15.00

<https://doi.org/10.1145/1122445.1122456>

## CCS CONCEPTS

• **Theory of computation** → **Data structures and algorithms for data management**; • **Information systems** → **Resource Description Framework (RDF)**; **Physical data models**.

## KEYWORDS

Hypertrie, RDF, SPARQL, triple store, worst-case optimal joins, graph patterns, graph database, hashing, path compression

## ACM Reference Format:

Anonymous Authors. 2018. Hashing the Hypertrie: Space- and Time-Efficient Indexing for SPARQL in Tensors. In *Woodstock '18: ACM Symposium on Neural Gaze Detection, June 03–05, 2018, Woodstock, NY*. ACM, New York, NY, USA, 10 pages. <https://doi.org/10.1145/1122445.1122456>

## 1 INTRODUCTION

The hypertrie [6], a monolithic indexing data structure based on tries, is designed to support efficiently evaluating SPARQL basic graph patterns (BGPs). While the access order for the positions of the tuples in tries is fixed, a hypertrie allows to iterate or resolve tuple positions in arbitrary order. In [6], Biegerl et al. show that hypertries of depth 3 are both time and memory efficient when combined with a worst-case optimal join (WCOJ) based on the Einstein summation algorithm. With the benchmarking of their implementation, dubbed TETRIS, they also show that hypertries outperform mainstream triple stores significantly on both synthetic and real-world benchmarks when combined with WCOJs.

We analyzed hypertries on four RDF datasets: Semantic Web Dog Food (SWDF), DBpedia 2015-10, WatDiv, and Wikidata (see Section 5 for details on the datasets). An analysis of the space requirements of the current implementation of hypertries revealed the following limitations:

- (1) The hypertries contain a high proportion of duplicate nodes. Up to 35% (Wikidata) of inner nodes and up to 91% (WatDiv) of leaf nodes were duplicates (see Figure 2).
- (2) Many nodes contain only a single entry. As shown in Figure 2, this applies to up to 62% of deduplicated inner nodes (DBpedia) and up to 75% of deduplicated leaf nodes (Wikidata).

Two main observations can be derived from this analysis. First, the duplicate nodes lead to an unnecessarily high memory footprint. The addition of deduplication to hypertries could hence yield an improved data structure with lower memory requirements. Second, the high number of single-entry nodes might lead to both unnecessary memory consumption and suboptimal query runtimes. A modification of the data structure to accommodate single-entry nodes effectively has the potential to improve both memory footprint and query runtimes.

Based on these findings, we propose the main contributions of this paper, i.e., the following three optimizations, which we illustrate in Figure 1:

- (1) **Hash-based identifiers (h):** We use hashes of nodes in hypertries as primary keys. Hence, nodes with the same entries are stored exactly once, thus eliminating duplicates.
- (2) **Single-entry node (s):** Single-entry nodes of height  $\geq 2$  (i.e., inner nodes) store the sub-hypertries of which they are the root node directly, thus saving space and eventually eliminating child nodes.
- (3) **In-place storage (i):** Boolean-valued height-1 single-entry nodes (leaf nodes) are eliminated completely. Their payload is stored in their parent nodes in the place where they otherwise would be referenced.

By applying all three techniques, the number of stored nodes is reduced by 82–90% (SWDF, WatDiv), and the memory consumption is reduced by 58–70% (SWDF, WatDiv), while the number of queries answered per second increases by up to four orders of magnitude on single queries.

The rest of this paper is structured as follows. First, we discuss related work in Section 2. In Section 3, we specify notations and conventions, introduce relevant concepts and describe the baseline hypertrie. We present our optimizations of the hypertrie in Section 4 and evaluate our optimized hypertries in Section 5. Finally, we conclude in Section 6.

## 2 RELATED WORK

The idea for single-entry node and in-place storage is based on path compression, a common technique to reduce the number of nodes required to encode a tree by storing non-branching paths in a single node. It was first introduced by Morrison in PATRICIA trees. [14]

Using hashing for deduplication, like in the proposed hypertrie context for hypertrie nodes (see Section 4.1.2 for details), is a well-known and widely used technique. [13]

Many query engines for RDF graphs have been proposed in recent years. [1, 3, 6, 9–11, 16, 17, 20, 22] Different engines deploy different mixes of indices and have different query execution approaches partly dependent on their indices. A common approach among SPARQL engines is to build multiple full indices in different

collation orders such as Fuseki [10], Virtuoso [9], Blazegraph [20], and GraphDB [17]. Some systems build additional partial indices on aggregates such as RDF-3X [16], or cache data for frequent joins such as gStore [22] for star joins. Building more indexes provides more flexibility in reordering joins to support faster query execution, while fewer indexes accelerate updates and require less memory.

When it comes to worst-case optimal joins (WCOJs) [5], classical indexing reaches its limits as indices for all collation orders are required. A system that takes this approach is Fuseki-LTJ [11], which implements the WCOJ algorithm Leapfrog TrieJoin (LTJ) [21] within a Fuseki triple store with indices in all collation orders. Recent works also propose optimized data structures that provide more concise indices with support for WCOJs. Qdag [15] provide support for WCOJs based on an extension of quad trees. Redundancy in the quad tree is reduced by implementing it as a directed acyclic graph (DAG) and reusing equivalent subtrees. A Circle [3] stores Burrows-Wheeler-transformed ID triples in bent wavelet trees along with an additional index to encode the triples of an RDF graph. Both Qdag and Circle are succinct data structures that must be built at once and do not support updates. In their evaluation of Circle, Arroyuelo et al. showed that Qdag and Circle are very space efficient, and that Fuseki-LTJ and Circle answer queries faster than state-of-the-art triple stores such as Virtuoso and Blazegraph with respect to average and median response times. The Qdag performed considerably worse in the query benchmarks than all other systems tested.

The hypertrie that we strive to optimize in this paper is, like the Qdag, internally represented as a DAG. As with the Qdag, the DAG nature of the hypertrie reduces the space requirement from factorial to exponential by the tuple length. The reduction is accomplished by eliminating duplicates among equal subtrees.

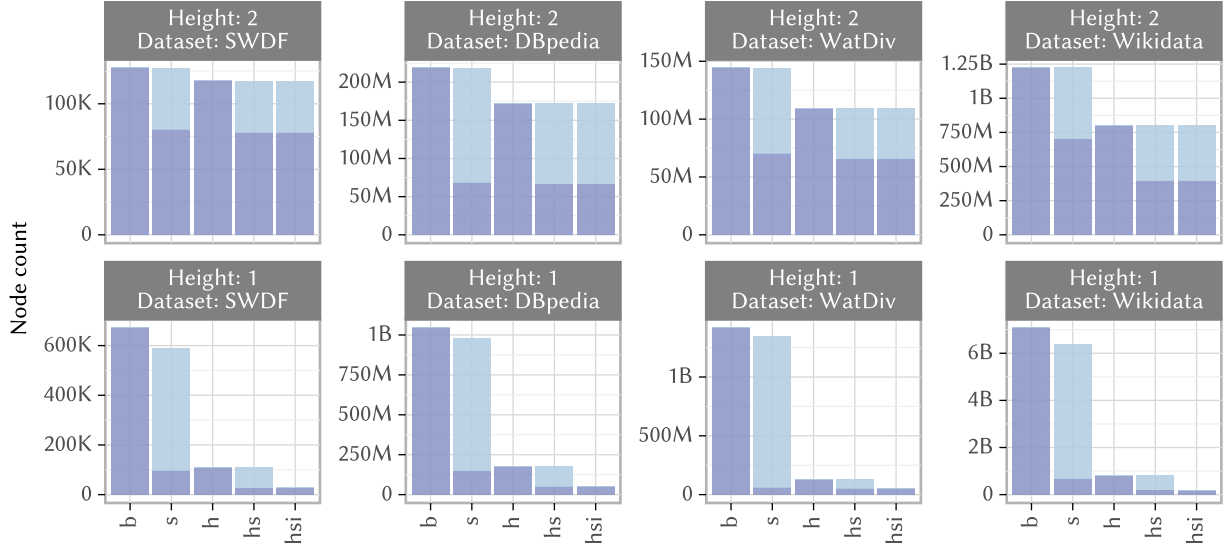
## 3 BACKGROUND

In this section, we briefly introduce the notation and conventions used in the rest of this paper. In particular, we give a brief overview of relevant aspects of RDF, SPARQL, and tensors. We also provide an overview of the formal specification of hypertries. More details can be found in [6].


### 3.1 Notation and Conventions

The conventions used in this paragraph stem from [6]. Let  $\mathbb{N}$  be the set of the natural numbers including 0. We use  $\mathbb{I}_n := \{i \in \mathbb{N} \mid 1 \leq i \leq n\}$  as a shorthand for the set of natural numbers from 1 to  $n$ . The domain of a function  $f$  is denoted  $\text{dom}(f)$  while  $\text{cod}(f)$  stands for the target (also called codomain) of  $f$ . A function which maps  $x_1$  to  $y_1$  and  $x_2$  to  $y_2$  is denoted by  $[x_1 \rightarrow y_1, x_2 \rightarrow y_2]$ . Sequences with a fixed order are delimited by angle brackets, e.g.,  $l = \langle a, b, c \rangle$ . Their elements are accessible via subscript, e.g.,  $l_1 = a$ . The number of times an element  $e$  is contained in any bag or sequence  $C$  is denoted by  $\text{count}(e, C)$ ; for example,  $\text{count}(a, \langle a, a, b, c \rangle) = 2$ . We denote the Cartesian product of  $S$  with itself  $i$  times with  $S^i = \underbrace{S \times S \times \dots \times S}_i$ .

We use the term *word* to describe a processor word, e.g. a 64-bit data chunk when using the x86-64 instruction set.



**Figure 2: Node counts of different hypertrie versions on four datasets. The hypertrie versions are identified by their features: baseline (b), single-entry node (s), hash identifiers (h), and in-place storage of height-1 single-entry nodes (i). The bars for baseline nodes are darker colored; the bars for single-entry nodes are lighter colored. Node counts are given for height-2 and height-1 nodes. The node count of height-3 nodes is always 1 baseline node for RDF tensors.**



**Figure 3: Example RDF graph. Integer indices for RDF resources are provided in parentheses behind the string identifier.**

### 3.2 RDF and SPARQL

An RDF statement is a triple  $\langle s, p, o \rangle$  and represents an edge  $s \xrightarrow{p} o$  in an RDF graph  $g$ .  $s$ ,  $p$  and  $o$  are called RDF resources. An RDF graph can be regarded as a set of RDF statements. The set of all resources of a graph  $g$  is given by  $r(g)$ . An example of an RDF graph is given in Figure 3. The graph contains, among others, the RDF statement  $\langle \text{:Alice}, \text{foaf:knows}, \text{:Bob} \rangle$ .

A triple pattern (TP)  $Q$  is a triple that has variables or RDF resources as entries, e.g.,  $\langle ?x, \text{foaf:knows}, ?y \rangle$ . Matching a triple pattern  $Q$  with a statement  $t$  results in a set of zero or one solution mappings. If  $Q$  and  $t$  have exactly the same resources in the same positions, then matching  $Q$  to  $t$  results in a solution mapping which maps the variables of  $Q$  to the terms of  $t$  in the same positions. For example, imagine  $Q = \langle ?x, \text{foaf:knows}, ?y \rangle$  and  $t = \langle \text{:Alice}, \text{foaf:knows}, \text{:Bob} \rangle$ . Then  $Q(t) = \{[?x \rightarrow \text{:Alice}, ?y \rightarrow \text{:Bob}]]\}$ .

Otherwise, the set of solutions is empty, i.e.,  $Q(t) = \emptyset$ . The result of matching a triple pattern  $Q$  against an RDF graph  $g$  is

$$Q(g) = \bigcup_{t \in g} Q(t), \quad (1)$$

i.e., the union of the matches of all triples  $t$  in  $g$  with  $Q$ . A list of triple patterns is called a basic graph pattern (BGP). The result of applying a BGP to an RDF graph  $g$  is the natural join of the solutions of its triple patterns.

Similar to previous works, previous works [4, 6, 16], we only consider the subset of SPARQL where a query is considered to consist of a BGP, a projection and a modifier (i.e., DISTINCT) that specifies whether the evaluation of  $Q$  follows bag or set semantics.

### 3.3 Tensors and RDF

Similar to [6], we use tensors that can be represented as finite multi-dimensional arrays. We consider a tensor of rank- $n$  as an  $n$ -dimensional array  $\mathbf{K}_1 \times \dots \times \mathbf{K}_n \rightarrow \mathbb{N}$  with  $\mathbf{K}_1 = \dots = \mathbf{K}_n \subset \mathbb{N}$ . Tuples from the tensor's co-domain  $\mathbf{k} \in \mathbf{K}$  are called keys. The entries  $k_1, \dots, k_n$  of a key  $\mathbf{k}$  are dubbed key parts. The array notation  $T[\mathbf{k}] = v$  is used to express that  $T$  stores for key  $\mathbf{k}$  value  $v$ .

The representation of  $g$  as a tensor, dubbed RDF tensor or adjacency tensor  $T$ , is a rank-3 tensor over  $\mathbb{N}$  which encodes  $g$ . Let  $id : r(g) \rightarrow \mathbb{I}_{|r(g)|+1}$  be an index function. The function maps each term of  $g$ —of which there are  $|r(g)|$ —to a fixed value in  $\mathbb{I}_{|r(g)|+1}$ . All unbound variables in solution mappings are mapped to  $|r(g)| + 1$ . For all statements  $\langle s, p, o \rangle \in g$ , the value of  $T[id(s), id(p), id(o)]$  is 1. All other values of  $T$  are 0. For example,  $T[5, 9, 7] = 1$  for the example graph shown in Figure 3, while  $T[5, 9, 6] = 0$ .

Matching a triple pattern  $Q$  against a graph  $g$  is equivalent to slicing the tensor representation of  $g$  with a slice key  $s(Q)$  corresponding to  $Q$ . The length of  $s(Q)$  is equal to the order of the tensor to which it is applied. Said slice key has a key part or a place holder, denoted “:” (no quotes), in every position. Slicing  $g$  with  $s(Q)$  results in a lower-order tensor that retains only entries where the key parts of the slice key match with the key parts of the tensor entries. For example, the slice key for the TP  $Q = \langle ?x, \text{foaf:knows}, ?y \rangle$  executed against the example graph in Figure 3 is  $\langle :, 8, : \rangle$ . Applying the TP  $Q$  to  $g$  is homomorphic to applying the slice key  $s(t)$  to the tensor representation of  $g$ .

To define a tensor representation for sets or bags of solutions, we first define an arbitrary but fixed ordering function *order* for variables (e.g., any alphanumeric ordering). A tensor representation  $T'$  of a set or bag of solutions is a tensor of rank equal to the number of projection variables in the query. The index for accessing entries of  $T'$  corresponds to *order*. For example, given the TP  $Q = \langle ?x, \text{foaf:knows}, ?y \rangle$  with the projection variables  $?x$  and  $?y$ ,  $T'$  would be a matrix with  $?x$  as the first dimension and  $?y$  as the second dimension. After applying  $Q$  to the graph in Figure 3, we would get a tensor  $T'$  with  $T'[2, 5] = 1$  and  $T'[5, 2] = 0$ .

The Einstein summation [8, 18] is an operation with variable arity. With this operation, the natural joins between the TPs of a BGP and variable projection can be combined into a single expression that takes the tensor representations of the TPs as input.

The execution of a SPARQL query on an RDF graph  $g$  is mapped to operations on tensors as follows. For each triple pattern, the RDF tensor  $T$  is sliced with the corresponding slice key. The slices are used as operands to an Einstein summation. Each slice is subscripted with the variables of the corresponding triple pattern. The result is subscripted with the projected variables. A ring with *addition* and *multiplication* is used to evaluate the Einstein summations. For example, evaluating the query with the BGP  $\langle \langle ?x, \text{foaf:knows}, ?y \rangle, \langle ?y, \text{rdf:type}, : \text{Pet} \rangle \rangle$  and a projection to  $?x$  on the RDF graph  $g$  from Figure 3 is equivalent to calculating  $\sum_x T[:, 8, :]_{x,y} \cdot T[:, 9, 7]_y$ , where  $T$  is the RDF graph of  $g$ .

### 3.4 Hypertrie

A hypertrie is a tensor data structure that maps strings of fixed length  $d$  over an alphabet  $A$  to some value space  $V$  [6]. It is implemented as a directed acyclic graph to store tensors sparsely by storing only non-zero entries. Formally, [6, p.62] defines a hypertrie as follows:

**DEFINITION 1 (HYPERTRIE).** Let  $H(d, A, E)$  with  $d \geq 0$  be the set of all hypertries with depth  $d$ , alphabet  $A$ , and values  $E$ . If  $A$  and  $E$  are clear from the context, we use  $H(d)$ . We set  $H(0) = E$  per definition. A hypertrie  $h \in H(1)$  has an associated partial function  $c_1^{(h)} : A \rightarrow E$  that specifies outgoing edges by mapping edge labels to children. For  $h' \in H(n)$ ,  $n > 1$ , partial functions  $c_p^{(h')} : A \rightarrow H(d-1)$ ,  $p \in \mathbb{I}_n$  are defined. Function  $c_p^{(h')}$  specifies the edges for resolving the part equivalent to depth  $p$  in a trie by mapping edge labels to children. For a hypertrie  $h$ ,  $z(h)$  is the size of the set or mapping it encodes.

An example of a hypertrie encoding the RDF tensor of the graph in Figure 3 is given in Figure 4 with the baseline hypertrie.

To retrieve the value for a tensor key, we start at the root node. If the current node is from  $H(0)$ , it is the value and we are done. Otherwise, we select a key part from the key at an arbitrary position  $p$ . If  $c_p$  maps the selected key part, we descend to mapped sub-hypertrie, remove the selected key part from the key and repeat the retrieval recursively on the sub-hypertrie with the shortened key. Otherwise, the value is 0.

Hypertries are designed to satisfy four conditions: (R1) memory efficiency, (R2) efficient slicing, (R3) slicing in any order of dimensions, and (R4) efficient iteration through slices. [6] Furthermore, note that every hypertrie is uniquely identified by the set of tuples it encodes.

*Implementation.* We refer to the original implementation of the hypertrie [6] as baseline implementation. The baseline hypertrie is implemented in C++. The lifetime of hypertrie nodes is managed by reference-counting memory pointers which free the memory of a node when it is no longer referenced. For nodes with height  $d > 1$ , the edge mappings  $c_p \in \mathbb{I}_d$  are stored in one hash table each. A node  $h'$  that is accessible from the root hypertrie  $h$  via different paths with equal slices is stored only once. Its parent nodes store a reference to the same physical instance of  $h'$ . For example, the slices  $h[3, :, :][:, 4]$  and  $h[:, :, 4][3, :]$  of a depth-3 hypertrie result in the same node. Nodes of depth  $d = 1$  store the leaf edges in a hash set.

Hypertries were introduced as a tensor data structure for the tensor-based triple store TETRIS. [6] In the following, we briefly describe the implementation of TETRIS, which is later used to evaluate the improvements to the hypertrie presented in this paper. Consider an RDF graph  $g$ . A depth-3 Boolean-valued hypertrie is used by TETRIS to store RDF triples encoded as integer triples. Therefore, the RDF resources  $r(g)$  are stored as heap-allocated strings. The integer identifier of a resource is its memory address. We write  $id(e)$  to denote the identifier of a resource  $e$ .  $id$  is implemented using a hash table while its inverse  $id^{-1}$  is applied by resolving the ID as memory address. Solutions of triple patterns are represented by pointers to sub-hypertrie nodes. Joins and projection are implemented with Einstein summation based on a worst-case optimal join algorithm.

## 4 APPROACH

In this section, we introduce three optimizations to the hypertrie. First, we eliminate duplicate nodes by identifying nodes with a hash. In a second step, we further reduce the memory footprint of hypertries by devising a more compact representation for nodes that encode only a single entry. Finally, we eliminate the separate storage of single-entry leaf nodes completely.

### 4.1 Hash-Based Identifiers

Our analysis of Figure 3 suggests that equal sub-hypertries are often stored multiple times. To eliminate this redundancy, we first introduce a hashing scheme for hypertries that can be updated incrementally. Based thereupon, we introduce the *hypertrie context*, which keeps track of existing hypertrie nodes and implements a hash-based deduplication.

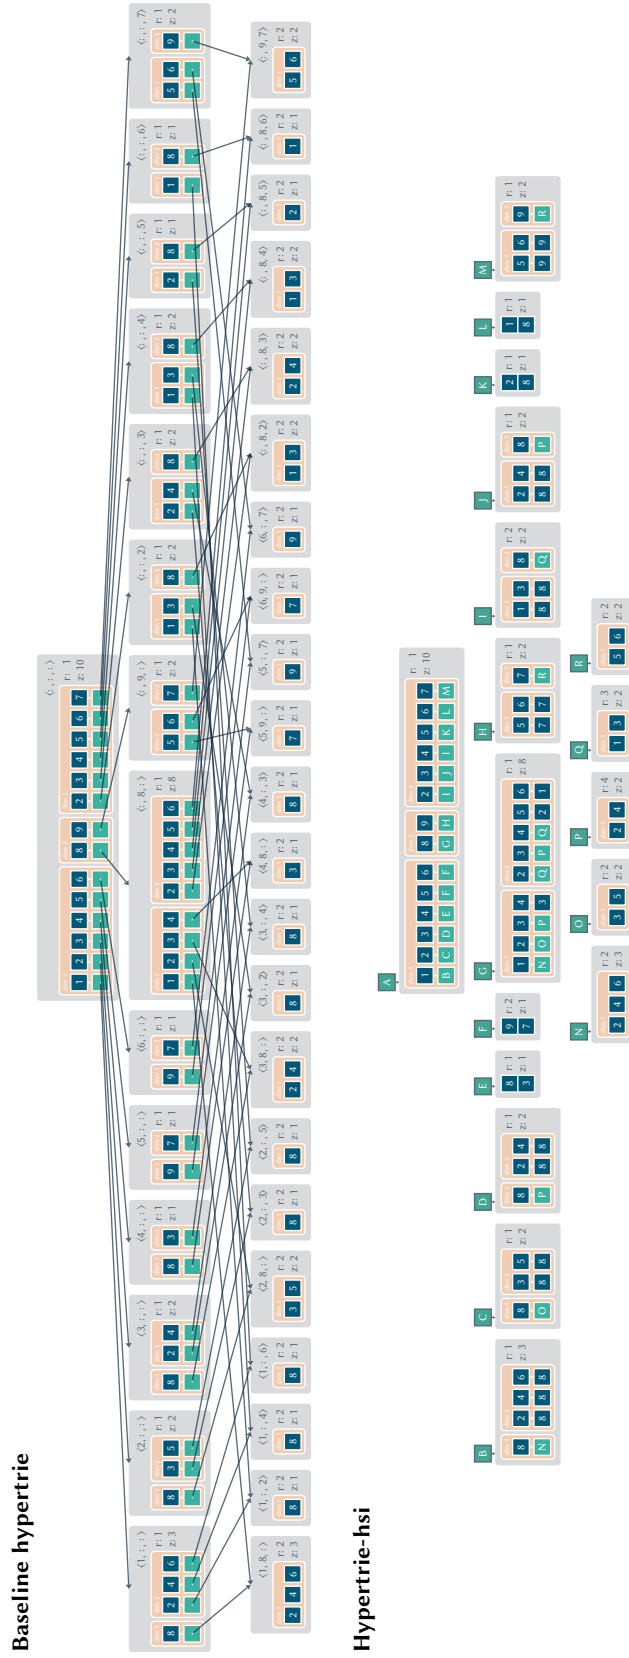


Figure 4: Top: baseline hypertrie; bottom: hypertrie-hsi with all proposed optimizations. Both hypertries encode the RDF graph from Figure 3. RDF resources are encoded by their ID (see also Figure 3).

**4.1.1 Hashing Hypertries.** To identify a suitable hash function, we exploit the fact that every hypertrie  $h$  is uniquely identified by the set of tuples it encodes. An efficient insertion (and deletion) of entries in  $h$  requires that a hash function for hypertries allows incremental updates in  $O(1)$  when adding (or removing) entries, rather than requiring to rehash all entries in  $O(z(h))$ . In the following, we define such a hashing scheme for a set of tuples.

Let  $i$  be a hash function that maps scalars to  $\mathbb{N}$ .<sup>1</sup> We extend  $i$  for lists by an order-dependent fold of the entries' hashes, i.e.,  $i(\langle a, b \rangle) \neq i(\langle b, a \rangle)$ . We further define  $i$  for sets by the fold of all entries' binary representations with XOR. An empty set is mapped to a default value, e.g., 0. Since XOR is self-inverse, commutative, and associative, the hash resulting from adding an element  $e$  to (or removing an element  $e$  from) a set  $s$  can be calculated incrementally by  $i(s) \oplus i(e)$ . Finally, we define the hash of a hypertrie as:

$$i(h) := i(\{k \mid k \in \text{dom}(h)\}).$$

**4.1.2 Hypertrie Context.** The goal of a hypertrie context is to ensure that hypertrie nodes are stored only once, regardless of how often they are referenced. We now describe the design requirements for hypertrie contexts, provide a formal definition, and conclude with implementation considerations.

In their baseline implementation, hypertrie nodes are retrievable by their path from the root of the hypertrie only. Information pertaining to the location of a node in memory is only available within its parent nodes. Consequently, only nodes with equivalent paths, i.e., with equal slice keys, are deduplicated in the baseline implementation. Equal hypertries with different slice keys are stored independently of each other. Hypertrie contexts eliminate these possible redundancies by storing hypertrie nodes by their hash and tracking how often nodes are referenced. The parent nodes are modified to reference their child nodes using hashes instead of memory pointers. Identifying hypertrie nodes by their hashes ensures that there are no duplicates.

A hypertrie can be *contained* or *primarily contained* in a hypertrie context  $hc$ . All nodes managed by a hypertrie context are *contained* therein. A hypertrie is said to be *primarily contained* in a hypertrie context  $hc$  iff it was stored explicitly in said context. For example, the root node of a hypertrie used for storing a given graph is commonly *primarily contained* in a hypertrie context. If a hypertrie  $h$  is *primarily contained* in a hypertrie context  $hc$ , then all sub-hypertries of  $h$  are *contained* in  $hc$ .

Adding a new primarily contained hypertrie or changing an existing hypertrie may alter the set of hypertries contained in a hypertrie context. To efficiently decide whether a node is still needed after a change, the hypertrie context tracks how often each node is referred to. Nodes that are no longer referenced after a change are removed. In *hypertrie contexts*, hypertries are considered to reference their sub-hypertries by hash.<sup>2</sup>

Formally, we define a *hypertrie context* as follows:

**DEFINITION 2 (HYPERTRIE CONTEXT).** Let  $A$  be an alphabet,  $E$  a set of values, and  $d \in \mathbb{N}$  the maximal depth of the hypertries that are to be stored.

We denote the set of hypertries  $\bigcup_{t \leq d} H(t, A, E)$  as  $\Lambda_0$ .  $\Lambda_0$  without empty hypertries  $\{h \in \Lambda_0 \mid z(h) \neq 0\}$  is denoted  $\Lambda$ .

A hypertrie context  $C$  for hypertries from  $\Lambda_0$  is defined by a triple  $(P, m, r)$  where

- $P$  is a bag of elements from  $\Lambda_0$ ,
- $m : \mathbb{Z} \rightarrow \Lambda$  maps hashes to non-empty hypertries which are  $P$  or are sub-hypertries of one of  $P$ 's elements, and
- $r : \Lambda \rightarrow \mathbb{N} \cup 0$  assigns a reference count to non-empty hypertries.

We define two relations between hypertrie context and hypertries:

- Hypertries  $p \in P$  are primarily contained in  $C$ , denoted as  $p \in C$ .
- Hypertries  $h \in \text{cod}(m)$  are contained in  $C$ , denoted as  $h \in C$ .

For a hypertrie  $h \in \Lambda$ ,  $r(h)$  is calculated from sum of the count of  $h$  in  $P$  and the number of references to  $h$  from hypertrie  $h' \in C$ :

$$r(h) := \text{count}(h, P) + \sum_{\substack{h' \in C \\ p \in \mathbb{I}_d}} \text{count}(h, \text{cod}(c_p^{(h')})).$$

## 4.2 Single-Entry Node

Central properties of a hypertrie are that slicing in any dimension can be carried out efficiently (see R2 and R3 in Section 3.4) and that non-zero slices can be iterated efficiently (see R4 in Section 3.4). In the implementation of hypertrie node described so far (in the following: full node), this is achieved by maintaining one hash table of non-zero slices for each dimension. The main observation behind this optimization is that R2–R4 also hold for a hypertrie node that represents only a single entry if the hypertrie node stores only the entry itself. We dub such a node *single-entry node* (SEN). A similar technique is used in radix trees [12] to store non-branching paths in a condensed fashion.

For slicing, it is sufficient to match the slice key against the single entry of the node. Thus, the result may have zero or one non-zero entry (see R2, R3). There is exactly one non-zero slice in each dimension. Iteration of the non-zero slices is now trivial (see R4).

SEN are—when applicable—always more memory efficient than full nodes.<sup>3</sup> Compared to a full node  $h$ , an SEN eliminates memory overhead in three ways. (1) It does not maintain hash tables  $c_p^{(h)}$  for edges to child nodes. (2) Child nodes do not need to be stored, unless they are also needed by other nodes. (3) The node size  $z(h)$  does not need to be stored explicitly since it is always 1.

Formally, we define an SEN as follows:

**DEFINITION 3 (SINGLE-ENTRY HYPERTRIE).** Let  $H$ ,  $d$ ,  $A$  and  $E$  be given as in Definition 1. Further, consider  $h \in H(d)$ , which stores for key  $\langle k_1, \dots, k_n \rangle$ , the value  $v$ . If  $h$  encodes exactly one entry ( $z(h) = 1$ ),  $h$  is defined as  $\langle \langle k_1, \dots, k_n \rangle, v \rangle$  and is called a single-entry node (SEN). Children mapping functions  $c_p^{(h)}$  are not defined for  $h$ .

<sup>1</sup>We used the hash functions from <https://github.com/martinus/robin-hood-hashing> since preliminary experiments showed that they performed well and had no collisions on the datasets from Section 5.

<sup>2</sup>The outgoing edges  $c_p^{(h)}$  in Definition 1 are considered to map hashes of hypertries instead of hypertries.

<sup>3</sup>Consider a hypertrie  $h$  with a single entry  $z(h) = 1$  and depth  $d \geq 1$ . A hash table that maps a key part requires more memory than just a single key part. Hence, the  $d$  child mapping hash tables of a full hypertrie node encoding  $h$  require more memory than the  $d$  key parts of the entry stored in an SEN encoding  $h$ . Consequently, an SEN node is always more memory efficient than a full node.

SEN can be used without limitations in a hypertrie context.

### 4.3 In-Place Storage

Our third optimization is to store certain nodes exactly where a reference to them would be stored otherwise. While the aforementioned optimizations can be used for hypertries with all value types (e.g., Boolean, integer, float), the optimization in this section is only applicable to Boolean-valued hypertries.

The payload of a binary-valued (note that our tensors only contain 0s and 1s) height-1 SEN is a single key part (1 word). It takes the same amount of memory as the hash that identifies the hypertrie (1 word) and which is stored in its parent nodes' children mappings to reference it. Therefore, the payload of a height-1 SEN fits into the place of its reference.

We use this property to reduce the total storage required: The payload of child height-1 SENs—their key part—is stored in place of their reference in the children mappings of their parent nodes. To encode if a hash or a key part is stored, a bit in the same fixed position of both key part and hash is reserved and used as a type tagging bit, e.g., the most significant bit. As in-place stored height-1 SEN are not heap-allocated, reference counting is not necessary. The memory is released properly when the hash table is destructed.

### 4.4 Example

An exemplary comparison of a baseline hypertrie and a hypertrie context containing one primary hypertrie with all three proposed optimizations is given in Figure 4.

## 5 EVALUATION

We implemented our optimizations within the TENTRIS framework. The goal of our evaluation was twofold: first, we assessed the index sizes and index generation times with four datasets of up to 5.5 B triples. In a second experiment, we evaluated the query performance of the triple stores in a stress test. Throughout our evaluation, we compared

- the original version of TENTRIS, dubbed TENTRIS-b,
- our extension of TENTRIS with hash identifiers (h) and single entry nodes (s), dubbed TENTRIS-hs,
- TENTRIS-hs extended with the in-place storage (i) optimization, dubbed TENTRIS-hsi, and
- the six popular triple stores, i.e., Blazegraph 2.1.6 Release Candidate, Fuseki 3.17.0, Fuseki-LTJ—a Fuseki that uses a worst-case optimal join algorithm—<sup>4</sup>, GraphDB 9.5.1, gStore 0.8<sup>5</sup>, and Virtuoso 7.2.5.1.

We chose popular triple stores which provide a standard HTTP SPARQL interface, support at least the same subset of SPARQL as TENTRIS and are freely available for benchmarking. We did not include Qdag or Ring because they do not provide a SPARQL HTTP endpoint and do not support projections. We used the datasets Semantic Web Dog Food (SWDF) (372 K triples), the English DBpedia version 2015-10 (681 M triples) and WatDiv [2] (1 B triples) and their respective query lists from [6]. We added Wikidata trusty from 2020-11-11 (5.5 B triples) as another large real-world dataset and

generated queries with FEASIBLE [19] from Wikidata query logs. As in [6], FEASIBLE was configured to generate SELECT queries with BGPs and DISTINCT as an optional solution modifier. All experiments were executed on a server with an AMD EPYC 7742, 1 TB RAM and two 3 TB NVMe SSDs in RAID 0 running Debian 10 and OpenJDK 11.0.11.

### 5.1 Index Size and Loading Time

Storage requirements for indices and index building speeds are reported in Figure 5. The index sizes of the TENTRIS versions were measured with `cgmemtime`'s<sup>6</sup> "Recursive and acc. high-water RSS+CACHE". For all other triple stores, the total size of the index files after loading was used. `cgmemtime`'s "Child wall" was used to measure the time for loading the datasets.

Three triple stores were not able to load the Wikidata dataset: gStore failed due to a limit on the number of used RDF Resources, Virtuoso failed because of a known bug<sup>7</sup>, and TENTRIS-b ran out of memory.

For all datasets, both optimized versions TENTRIS-hs and TENTRIS-hsi improve the storage efficiency of hypertries. Compared to TENTRIS-b, TENTRIS-hs has a 55–68% reduced storage requirement for the datasets SWDF, DBpedia, and WatDiv at the cost of 9–36% longer index building times, while TENTRIS-hsi has a 58–70% reduced storage requirement and takes 2–28% longer to build the indices. For the Wikidata dataset, the index sizes of TENTRIS-hs and TENTRIS-hsi are reduced by at least 39% and 42%<sup>8</sup>, respectively. The in-place storage of single-entry leaf nodes (i) in TENTRIS-hsi saves memory, (1–7%) compared to TENTRIS-hs, and speeds up the index building (2–57%) on all datasets. For the small to medium-sized datasets SWDF, DBpedia, and WatDiv, the index building is slightly faster by 2–7%; for the large dataset, Wikidata, the margin is considerably larger with 56% improvement.

The index sizes of all TENTRIS versions scale similarly to other triple stores. The TENTRIS-hsi indices are similar in size to the indices produced by other triple stores. Compared to the smallest index for each dataset, TENTRIS-hsi uses 1.14 to 4.24 times more space. The loading time of TENTRIS-hsi is close to the mean of the non-TENTRIS triple stores.

### 5.2 Querying Stress Test

Our evaluation setup was similar to that used in [6]. The experiments were executed using the benchmark execution framework IGUANA v3.2.1 [7]. For each benchmark, the query mix was executed 30 times on each triple store and the timeout for a single query execution was set to 3 minutes. We report the performance using Queries per Second (QpS), Query Mixes per Hour (QMpH) and the proportion of failed queries. For QpS, only query executions that were successful and finished before the timeout are considered. The reported QpS value of a query on a dataset and triple store is the mean of the single measurements. Failed queries are penalized with the timeout duration for QMpH. We chose to report both QMpH and QpS to get a more fine-grained view of the performance.

<sup>4</sup>Fuseki-LTJ is based on Apache Jena Fuseki 3.9.0


<sup>5</sup>We used the modified version from [6] that fixes the SPARQL endpoint and sets a query timeout.

<sup>6</sup><https://github.com/gsaathof/cgmemtime>

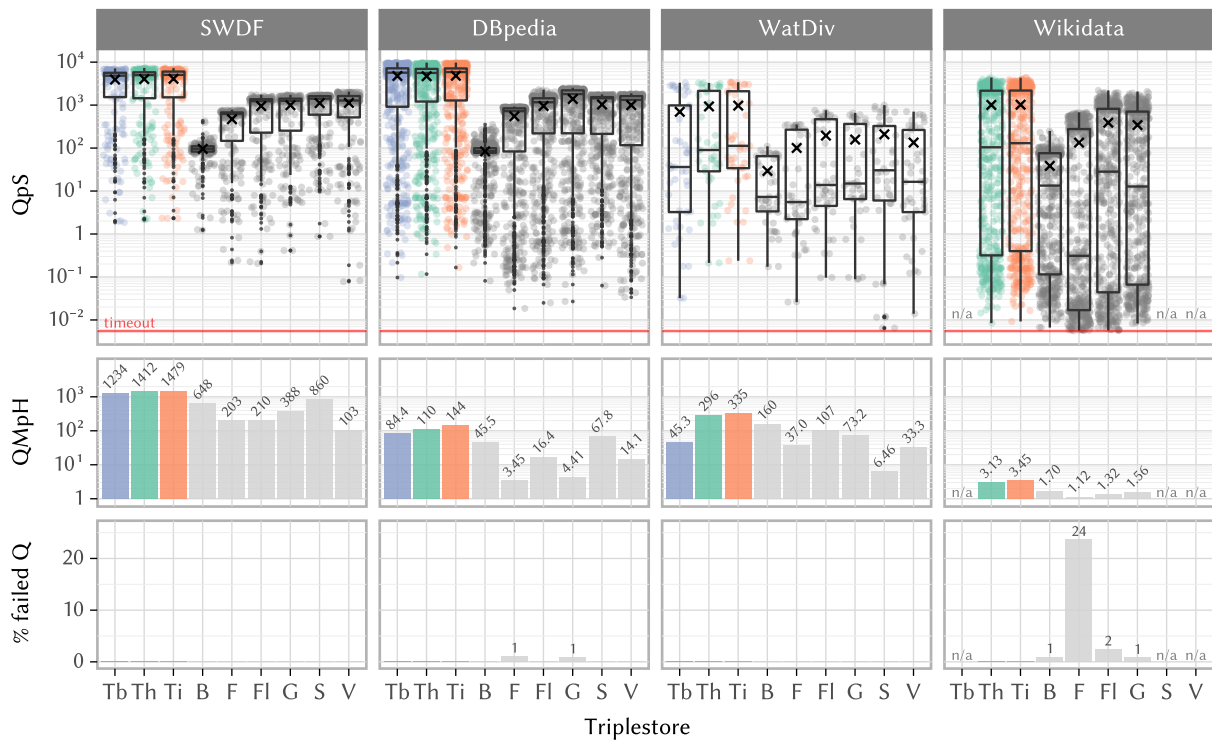
<sup>7</sup><https://community.openlinksw.com/t/loading-full-wikidata-latest-ttl-dump-into-vos/1880>, visited on Sept 14, 2021

<sup>8</sup>In comparison to 1TB RAM because TENTRIS-b ran out memory during loading.






**Figure 5: Storage efficiency and loading speed of TETRIS, with different hypertrie versions, and of other triple stores on four datasets. If loading a dataset with a triple store failed, the plot says n/a.**



**Figure 6: Performance metrics of query stress tests on four benchmarks for TETRIS, with different hypertrie versions, and for other triple stores. The first row of plots shows a boxplot and a scatterplot of Queries per Second (QpS). For this plot, only query executions that were successful and finished before the timeout are considered. Repeated measurements of a single query on a dataset and triple store are aggregated by mean into a single point in the scatterplot. The boxes indicate the first quartile, median, and third quartile. The whiskers have a length of 1.5 times the interquartile range. Outliers are marked with black dots, and crosses mark the means. The second row of plots shows Query Mixes per Hour (QmPH). Failed queries are rated with a timeout duration of 180 s. The third row of plots shows the percentage of queries that failed. If no number is provided, no queries failed. If experiments were not executed for a combination of triple store and benchmark, the plot says n/a. The triple stores benchmarked are TETRIS baseline (Tb), TETRIS-hs (Th), TETRIS-hsi (Ti), Blazegraph (B), Fuseki (F), Fuseki-LTJ (FI), GraphDB (G), gStore (S), and Virtuoso (V).**





**Figure 7: Heat map of Queries per Second (QpS) relative to TENTRIS-hsi (Ti) for each single query of the WatDiv benchmark. For triple store abbreviations, see Figure 6. gStore returned a wrong result on query C1. The corresponding cell is colored black.**

While QpS is more robust against outliers, QMpH can be strongly influenced by long-running and failed queries.

The baseline version of TENTRIS, TENTRIS-b, already performs better with respect to QpS and QMpH compared to all non-TENTRIS triple stores on the SWDF and DBpedia benchmarks but not on the WatDiv benchmark. Here, gStore and Blazegraph outperform TENTRIS-b by a factor of 1.6 and 3.5 with respect to QMpH.

TENTRIS-hs and TENTRIS-hsi both outperform TENTRIS-b and all other triple stores on all datasets with respect to average QpS (avgQpS) and QMpH. For the small real-world dataset SWDF, all TENTRIS versions answered queries with similar avgQpS ranging from 3935 (TENTRIS-b) to 4088 (TENTRIS-hsi, +4%). The same holds true for the larger real-world dataset DBpedia, with avgQpS ranging from 4753 (TENTRIS-b) to 4825 (TENTRIS-hsi, +1.5%). With respect to QMpH, the optimized TENTRIS versions clearly outperform TENTRIS-b by 14% (TENTRIS-hs) and 20% (TENTRIS-hsi) on the SWDF dataset, and even by 31% and 71% on the DBpedia dataset. On the synthetic dataset, the optimized TENTRIS versions show notable speedups on both metrics, avgQpS and QMpH. AvgQpS is increased from 698 by TENTRIS-b to 931 (+33%) by TENTRIS-hs and to 972 (+39%) by TENTRIS-hsi. QMpH is 6.6 and 7.4 times higher with TENTRIS-hs and TENTRIS-hsi, respectively, than with TENTRIS-b.

On Wikidata, measurements are available only for the TENTRIS versions hs and hsi due to TENTRIS-b not being able to load the dataset. TENTRIS-hsi is again slightly faster than TENTRIS-hs, with 1009 (hs) and 1021 (+1%, hsi) avgQpS, and 3.13 (hs) and 3.45 (+9%, hsi) QMpH. When compared to the fastest non-TENTRIS triple store on each metric and dataset, TENTRIS-hsi is 3–3.7 times faster with respect to avgQpS and 1.7–2.1 times faster with respect to QMpH.

None of the TENTRIS versions had failed queries during execution. On the DBpedia dataset, Fuseki and gStore failed on about 1% of the queries. On the Wikidata dataset, all non-TENTRIS triples stores that succeeded to load the dataset failed on some queries.

For the synthetic WatDiv benchmark, information on the shapes of queries is available. The queries are classified as linear queries

(L), star queries (S), snowflake-shaped queries (F), complex queries (C), incremental linear queries (IL) or mixed linear queries (ML). Queries C1-D and C2-D use set semantics, all other queries use bag semantics. For each query and triple store, the performance with respect to QpS relative to TENTRIS-hsi is shown in Figure 7. The optimizations show a notable effect on most star queries and on some snowflake-shaped and complex queries. Star queries S2, S4, and S6 are answered about two orders of magnitude faster by TENTRIS-hs than by TENTRIS-hsi. S1 is answered even three orders of magnitude faster. The snowflake-shaped queries F1 and F2 are answered three orders of magnitude faster by TENTRIS-hsi than by TENTRIS-b. The same speed-up is reached for the complex query C1 with both set and bag semantics.

With respect to QpS on the WatDiv benchmark, TENTRIS-hs is for every query about equal fast as TENTRIS-b or faster. The same applies to TENTRIS-hsi and TENTRIS-hs. Compared to all other triple stores, there are only three queries where another triple store is faster than TENTRIS-hsi. Virtuoso is about 20% faster than on query C2 with bag semantics, and gStore is on queries S5 and F3 about 2.1 and 2.5 times faster.

### 5.3 Discussion

The evaluation shows that applying all three optimizations (hsi) is in all aspects superior to applying only the first two optimizations (hs). Thus, we will consider only TENTRIS-hsi in the following. The proposed optimizations of the hypertrie improve the storage efficiency by 70% and the query performance with respect to avgQpS by large margins of up to four orders of magnitude. These improvements come at the cost of slightly longer index building times of at most 28%. The optimization of the storage efficiency is clearly attributable to the reduced number of nodes, as shown in Figure 2. For the improved query performance, definite attribution is difficult. We worked out two main factors we believe are reasonable to assume as the cause: First, information that was stored in a node and its subnodes in the baseline version is in the optimized version

more often stored in a single node. This way, the optimizations single-entry node (s) and in-place storage (i) cause fewer CPU cache misses and fewer resolves of memory addresses, resulting in faster execution. Second, key parts are not stored necessarily in a hash table anymore. Whenever a key part is read from a single-entry node (s) or in-place stored node (i), the optimized version saves one hash table lookup compared to the baseline version. On the other side, additional hash table lookups are required to retrieve nodes by their hash identifiers during query evaluation. We minimize this overhead by handling nodes by their memory address during evaluation after they were looked up by their hash first. The memory overhead for storing these handles is negligible as typically only a few are required at the same time.

For triple stores, there is always a trade-off between storage efficiency, index build time, and query performance. For example, less compressed indices can typically be built faster. Building multiple indices takes longer, but multiple indices allow for more optimized query plans. The baseline hypertrie clearly attributed significant weight to good query performance, with average index building time and above average storage requirement. The optimized hypertrie trades only a little worse index building time for better query performance and much-improved storage efficiency. The result is a triple store with superior query performance, average storage requirement, and still average index building time. Given the predominantly positive changes in trade-offs, we consider the proposed optimizations a substantial improvement.

## 6 CONCLUSION AND OUTLOOK

We presented a memory-optimized version of the hypertrie data structure. The three optimizations of hypertries that we developed and evaluated improved both the memory footprint and query performance of hypertries. A clear but small trade-off of our approaches is the slightly longer index building time they require.

The new storage scheme for hypertrie opens up several new avenues for future improvements. The persistence of optimized hypertrie nodes is easier to achieve due to the switch from memory pointers to hashes. Furthermore, the hash identifiable hypertrie nodes provide the building bricks to distribute a hypertrie over multiple nodes in a network. For TETRIS, the introduction of the hypertrie context opens up the possibility to store the hypertries of multiple RDF graphs in a single context and thereby automatically deduplicate common sub-hypertries. Especially for similar graphs, this optimization has the potential to improve storage efficiency substantially.

## REFERENCES

- [1] Waqas Ali, Muhammad Saleem, Bin Yao, Aidan Hogan, and Axel-Cyrille Ngonga Ngomo. 2021. A Survey of RDF Stores & SPARQL Engines for Querying Knowledge Graphs. arXiv:2102.13027 [cs.DB]
- [2] Güneş Aluç, Olaf Hartig, M. Tamer Özsu, and Khuzaima Daudjee. 2014. Diversified Stress Testing of RDF Data Management Systems. In *The Semantic Web – ISWC 2014*. Springer International Publishing, Cham, 197–212.
- [3] Diego Arroyuelo, Aidan Hogan, Gonzalo Navarro, Juan L. Reutter, Javiel Rojas-Ledesma, and Adrián Soto. 2021. *Worst-Case Optimal Graph Joins in Almost No Space*. Association for Computing Machinery, New York, NY, USA, 102–114. <https://doi.org/10.1145/3448016.3457256>
- [4] Medha Atre, Vineet Chaoji, Mohammed J. Zaki, and James A. Hendler. 2010. Matrix “Bit” Loaded: A Scalable Lightweight Join Query Processor for RDF Data. In *Proceedings of the 19th International Conference on World Wide Web (Raleigh, North Carolina, USA) (WWW ’10)*. Association for Computing Machinery, New York, NY, USA, 41–50. <https://doi.org/10.1145/1772690.1772696>
- [5] Albert Atserias, Martin Grohe, and Daniel Marx. 2008. Size Bounds and Query Plans for Relational Joins. In *49th Annual IEEE Symposium on Foundations of Computer Science, FOCS 2008, October 25–28, 2008, Philadelphia, PA, USA*. IEEE Computer Society, 739–748. <https://doi.org/10.1109/FOCS.2008.43>
- [6] Alexander Biggerl, Felix Conrads, Charlotte Behning, Mohamed Ahmed Sherif, Muhammad Saleem, and Axel-Cyrille Ngonga Ngomo. 2020. TETRIS – A Tensor-Based Triple Store. In *The Semantic Web – ISWC 2020*. Springer International Publishing, Cham, 56–73.
- [7] Felix Conrads, Jens Lehmann, Muhammad Saleem, Mohamed Morsey, and Axel-Cyrille Ngonga Ngomo. 2017. Iguana: A Generic Framework for Benchmarking the Read-Write Performance of Triple Stores. In *The Semantic Web – ISWC 2017*. Springer International Publishing, Cham, 48–65.
- [8] A. Einstein. 1916. Die Grundlage der allgemeinen Relativitätstheorie. *Annalen der Physik* 354 (1916), 769–822. <https://doi.org/10.1002/andp.19163540702>
- [9] Orri Erling. [n.d.]. Virtuoso, a Hybrid RDBMS/Graph Column Store. <http://vos.openlinksw.com/owiki/wiki/VOS/VOSArticleVirtuosoAHybridRDBMSGraphColumnStore> Retrieved Mar 17, 2018 from <http://vos.openlinksw.com/owiki/wiki/VOS/VOSArticleVirtuosoAHybridRDBMSGraphColumnStore>
- [10] Apache Software Foundation. 2019. Apache Jena Documentation - TDB Architecture. <https://jena.apache.org/documentation/tdb/architecture> Retrieved Apr 25, 2019 from <https://jena.apache.org/documentation/tdb/architecture.html>
- [11] Aidan Hogan, Cristian Riveros, Carlos Rojas, and Adrián Soto. 2019. A Worst-Case Optimal Join Algorithm for SPARQL. In *The Semantic Web – ISWC 2019 (Lecture Notes in Computer Science)*. Springer International Publishing, Cham, 258–275. [https://doi.org/10.1007/978-3-030-30793-6\\_15](https://doi.org/10.1007/978-3-030-30793-6_15)
- [12] V. Leis, A. Kemper, and T. Neumann. 2013. The adaptive radix tree: ARTful indexing for main-memory databases. In *2013 IEEE 29th International Conference on Data Engineering (ICDE)*. 38–49. <https://doi.org/10.1109/ICDE.2013.6544812>
- [13] Jyoti Malhotra and Jagdish Bakal. 2015. A survey and comparative study of data deduplication techniques. In *2015 International Conference on Pervasive Computing (ICPC)*. 1–5. <https://doi.org/10.1109/PERVASIVE.2015.7087116>
- [14] Donald R. Morrison. 1968. PATRICIA—Practical Algorithm To Retrieve Information Coded in Alphanumeric. *J. ACM* 15, 4 (Oct. 1968), 514–534. <https://doi.org/10.1145/321479.321481>
- [15] Gonzalo Navarro, Juan L. Reutter, and Javiel Rojas-Ledesma. 2020. Optimal Joins Using Compact Data Structures. In *23rd International Conference on Database Theory, ICDT 2020, March 30–April 2, 2020, Copenhagen, Denmark (LIPIcs, Vol. 155)*. Schloss Dagstuhl - Leibniz-Zentrum für Informatik, 21:1–21:21. <https://doi.org/10.4230/LIPIcs.ICDT.2020.21>
- [16] Thomas Neumann and Gerhard Weikum. 2008. RDF-3X: A RISC-Style Engine for RDF. *Proc. VLDB Endow.* 1, 1 (Aug. 2008), 647–659. <https://doi.org/10.14778/1453856.1453927>
- [17] Inc. Ontotext USA. [n.d.]. Storage — GraphDB Free 8.9 documentation. <http://graphdb.ontotext.com/documentation/free/storage.html#storage-literal-index> Retrieved Apr 16, 2019 from <http://graphdb.ontotext.com/documentation/free/storage.html#storage-literal-index>
- [18] MMG Ricci and Tullio Levi-Civita. 1900. Méthodes de calcul différentiel absolu et leurs applications. *Math. Ann.* 54, 1-2 (1900), 125–201.
- [19] Muhammad Saleem, Qaiser Mehmood, and Axel-Cyrille Ngonga Ngomo. 2015. FEASIBLE: A Feature-Based SPARQL Benchmark Generation Framework. In *The Semantic Web – ISWC 2015*. Springer International Publishing, Cham, 52–69.
- [20] SYSTAP, LLC. 2013. Bigdata Database Architecture - Blazegraph. [https://blazegraph.com/docs/bigdata\\_architecture\\_whitepaper.pdf](https://blazegraph.com/docs/bigdata_architecture_whitepaper.pdf) Retrieved Nov 29, 2019 from [https://blazegraph.com/docs/bigdata\\_architecture\\_whitepaper.pdf](https://blazegraph.com/docs/bigdata_architecture_whitepaper.pdf)
- [21] Todd L. Veldhuizen. 2014. Triejoin: A Simple, Worst-Case Optimal Join Algorithm. In *Proc. 17th International Conference on Database Theory (ICDT), Athens, Greece, March 24–28, 2014*. OpenProceedings.org, 96–106. <https://doi.org/10.5441/002/icdt.2014.13>
- [22] Lei Zou, M. Tamer Özsu, Lei Chen, Xuchuan Shen, Ruizhe Huang, and Dongyan Zhao. 2014. GStore: A Graph-Based SPARQL Query Engine. *The VLDB Journal* 23, 4 (2014), 565–590. <https://doi.org/10.1007/s00778-013-0337-7>